

Portfolio

By: Jorge Victor Turriate
Llallire

Summary

1. Introduction
2. Master's thesis
3. Projects

INTRODUCTION

Profile

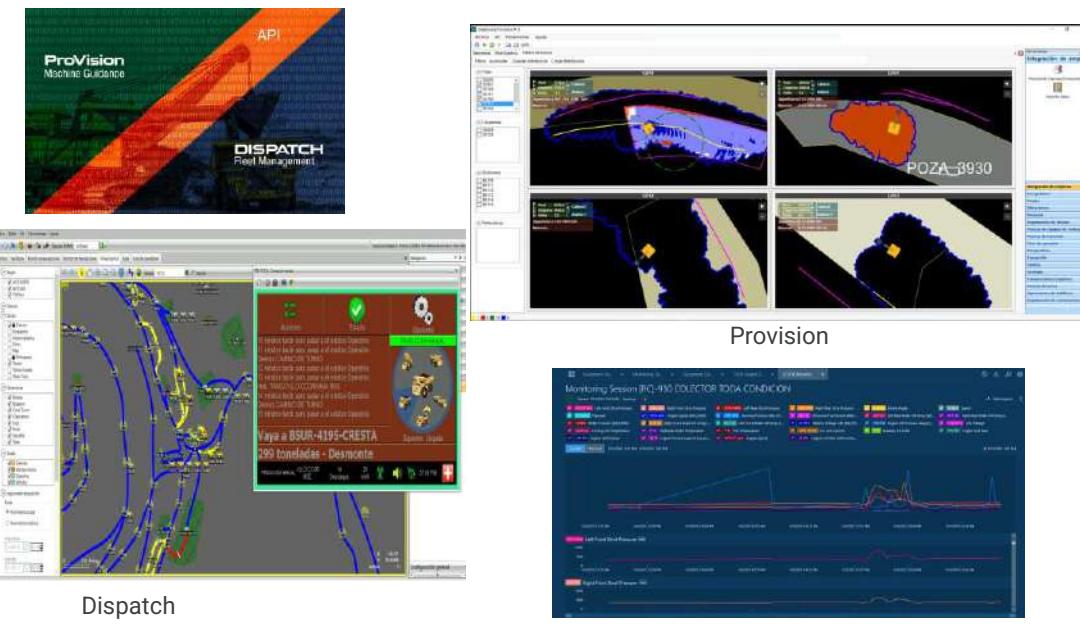
AI Researcher,
Telecommunications
Engineer and Web
developer

Education

Bachelor in Telecommunications Engineering



Master of Science in Artificial Intelligence and Computer Vision



Modular Mining Products

Professional Experience

Web developer

August 2019 – November 2019



Project Engineer Trainee

December 2019 – April 2021

Software, Deployment and Support Engineer

May 2021 – September 2023



FrontRunner AHS



Project leadership



Troubleshooting

MASTER'S THESIS

Curriculum Learning for Monocular Depth Estimation (MDE) in UAV Disaster Management

CONTEXT:

Disaster management:

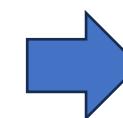
Requires rapid and accurate decision-making



Crowd counting:
Identify regions of interest



Depth Estimation:
Help assess terrain and generate depth maps



Monocular Depth Estimation (MDE)

Allows depth prediction from a single image

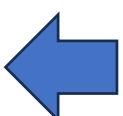


MDE Problem:

Training MDE models is computationally intensive and often requires large, diverse datasets and a lot of training hours.

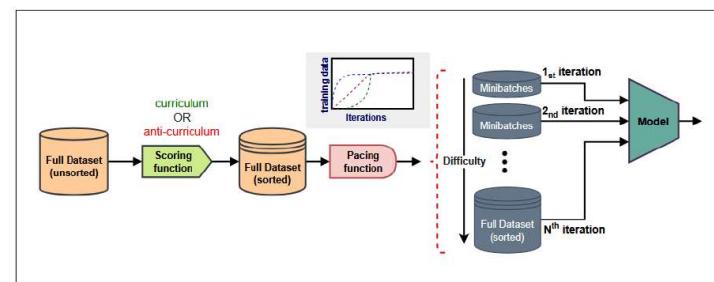


Will Curriculum Learning help the training of MDE models?



Curriculum Learning:

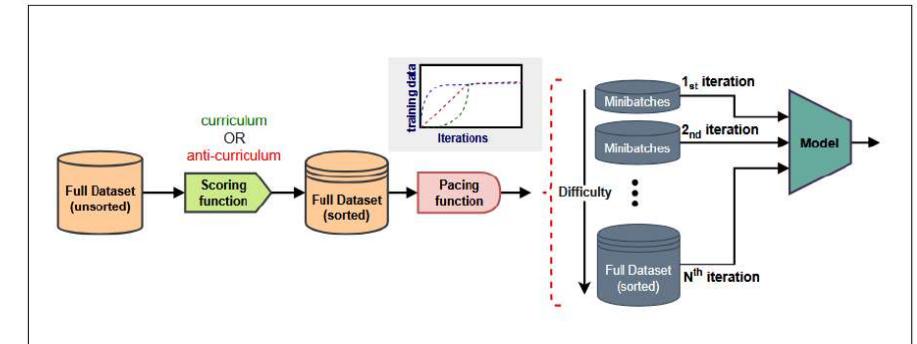
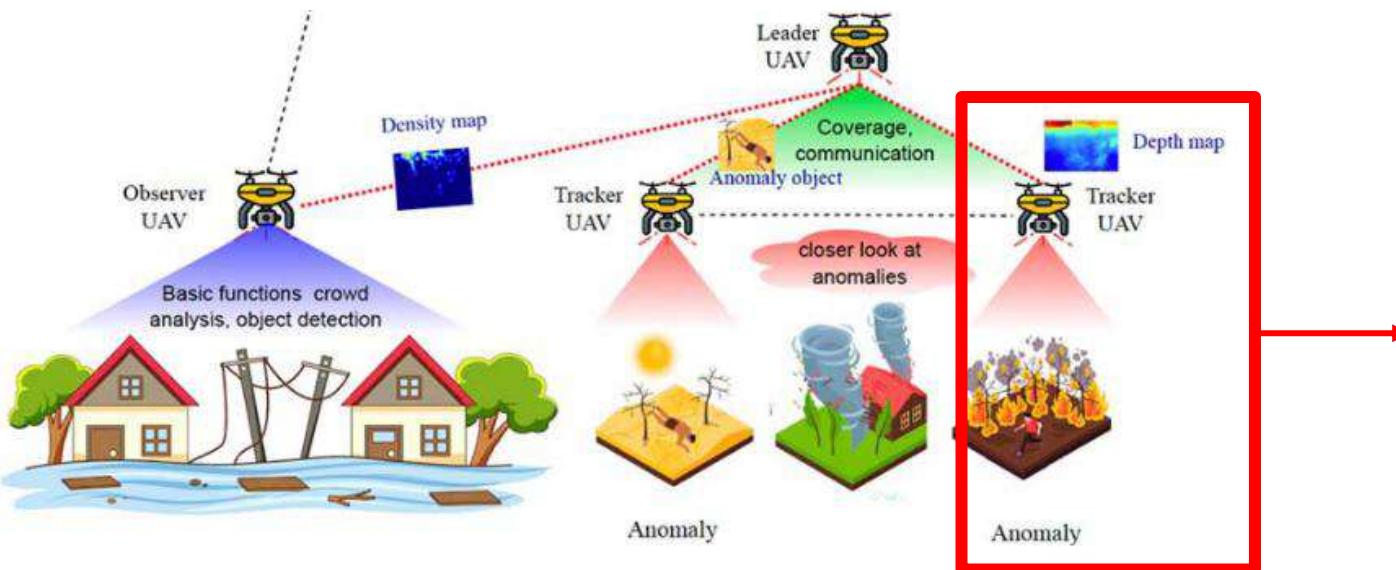
Offers a potential solution by accelerating convergence and improving generalization by mimicking human learning by presenting training data in a meaningful sequence



MISSION: Investigate the feasibility and impact of applying Curriculum Learning to Monocular Depth Estimation for UAVs used in disaster management scenarios.

OBJECTIVES:

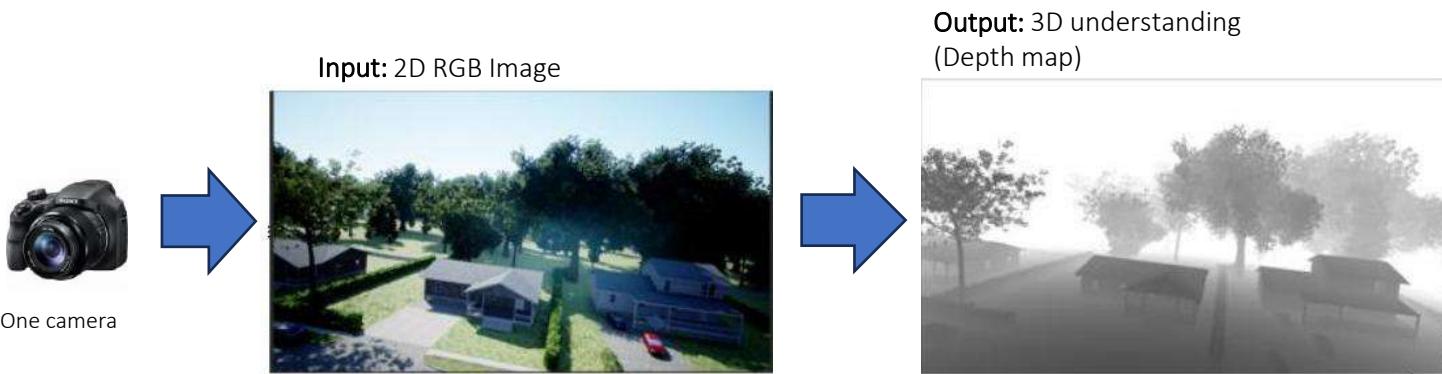
- Review the current state-of-the-art MDE models suitable for UAVs.
- Adapt and structure training datasets to incorporate curriculum strategies.
- Compare training performance and accuracy between curriculum-based training and traditional training.



MONOCULAR DEPTH ESTIMATION (MDE)

Definition:

- Predict Depth map from a single RGB image
- Key Visual cues: Perspective, shading, texture, occlusion



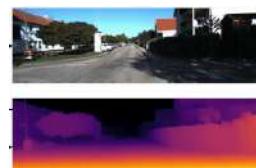
Comparison with other Depth Estimation models:

- Stereo Vision
- RGB-D
- Active sensors (LIDAR)

MDE has a **low cost**

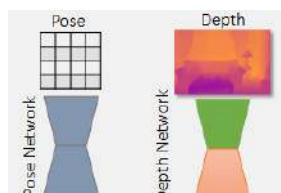
Deep Learning paradigms for MDE:

- Supervised Learning



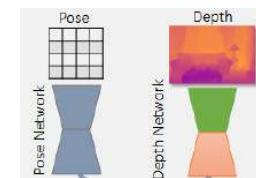
Training using the GT
Depth map

- Unsupervised Learning (Self-supervised Learning)

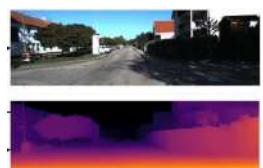


Training without any GT

- Hybrid approach



Pretrained network (SSL)



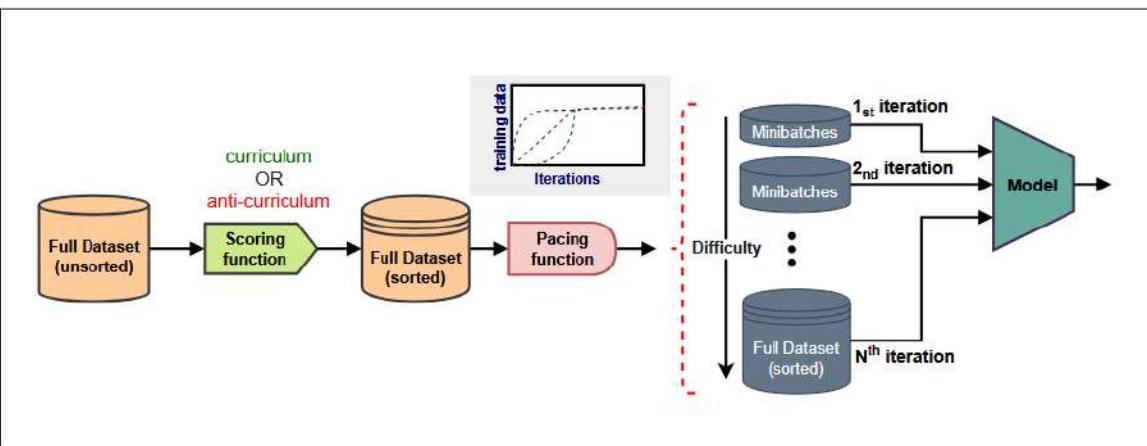
Fine-tuning
(Supervised learning)

CURRICULUM LEARNING (CL)

Definition:

Training paradigm inspired by natural learning in humans.

Core principle: Structure the data in progressive manner based on the order.



Algorithm 1: Curriculum learning

Input: T steps, N samples, Initial weights w^0 , training set $X = \{x_1, x_2, \dots, x_N\}$, pacing function $g: [T] \rightarrow [N]$, scoring function $s: [N] \rightarrow \mathbb{R}$, order $o \in \{\text{"ascending"}, \text{"descending"}\}$

```

 $(x_1, x_2, \dots, x_N) \leftarrow \text{sort } X \text{ using } s(x) \text{ in order } o$ 
for  $t = 1, 2, 3, \dots, T$  do
|  $\text{size} \leftarrow g(t)$ 
|  $X_t = X[1:\text{size}]$ 
|  $w^{(t)} \leftarrow \text{train\_one\_epoch}(w^{(t-1)}, X_t)$  sampling uniformly mini batches
end for
  
```

Elements:

- **Scoring function**

Defines **the score** for a given training sample

The **difficulty** is defined with the comparison of scores

$$s(x_j, y_j) > s(x_i, y_i)$$

The **loss** is used to define the scores

$$s(x_i, y_i) = \text{loss}(f_w(x_i), y_i)$$

- **Order:**

- ✓ Curriculum: From lowest to highest score
- ✓ Anti-Curriculum: From highest to lowest score

Types of scoring function:

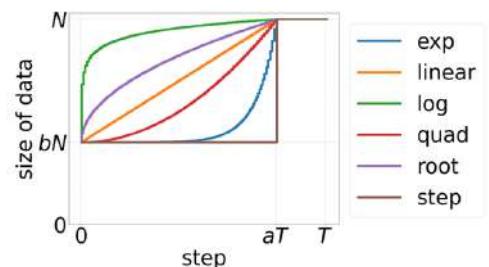
- ✓ Self-taught scoring function
- ✓ Transfer scoring function

- **Pacing function**

Determines **the size** of the training data subset to be used at iteration T

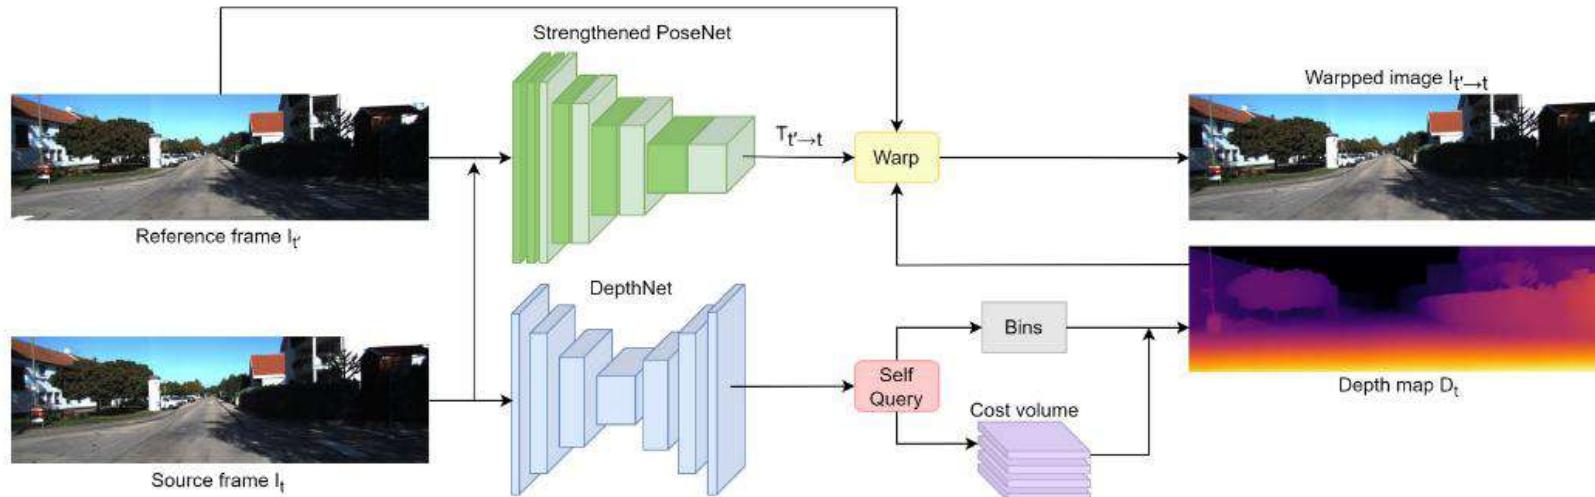
- a : Defines the speed in which the pacing function reaches the total number of samples
- b : fraction of training data used at the beginning

Name	Expression $g_{(a,b)}(t)$
linear	$Nb + N \frac{(1-b)}{(aT)} t$
quadratic	$Nb + N \frac{(1-b)}{(aT)^2} t^2$
root	$Nb + N \frac{(1-b)}{(aT)^{1/2}} t^{1/2}$
logarithmic	$Nb + N(1-b)(1 + 0.1 \log(\frac{t}{aT}) + e^{-10})$
exponential	$Nb + \frac{N(1-b)}{e^{10}-1} (\exp(\frac{10t}{aT}) - 1)$
step	$Nb + N[\frac{x}{aT}]$



SPId depth network

SOTA MDE network. It uses **self-supervised learning** to train its subnetworks.



Sub Networks:

- **Depth Net:** Infers depth maps from single RGB images by using an encoder-decoder framework.
- **Pose Net:** Predicts the relative pose (T') between the image I and the reference image I'

Loss function:

$$L = \mu L_p + \lambda L_s$$

- L_p : Masked photometric loss
- L_s : Per-pixel smooth loss

Training strategy:

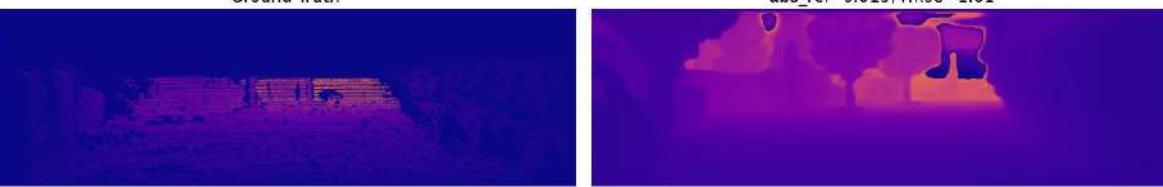
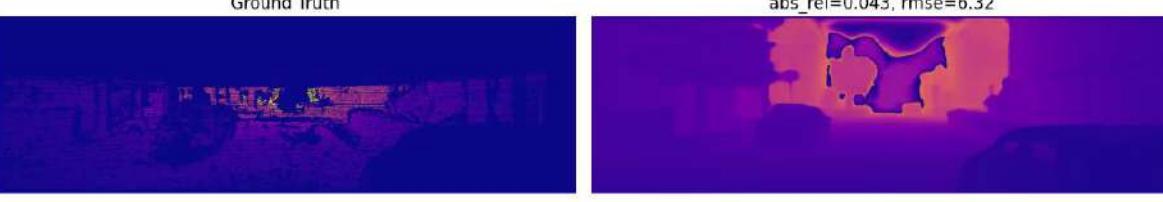
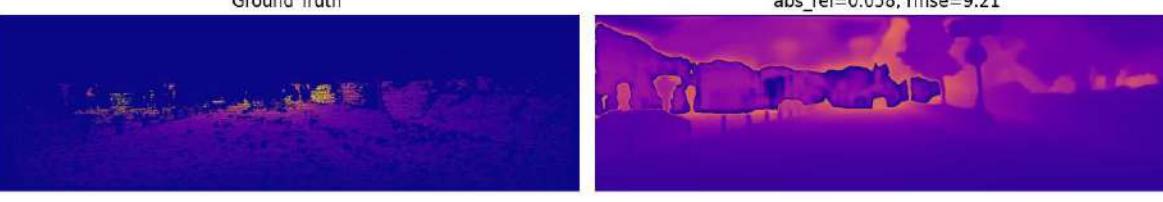
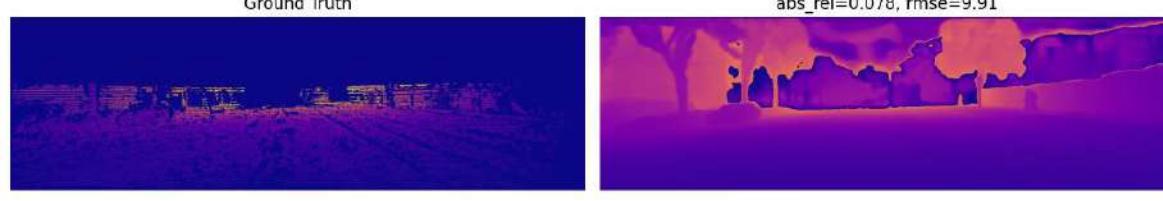
It uses a **hybrid learning** approach

- **Self-supervised training** on 75,000 images from KITTI
- **Supervised fine-tuning** on additional 25,000 images from KITTI

Inference on KITTI Eigen Split:

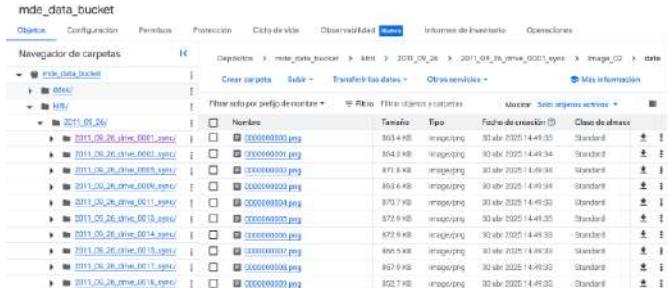
Absolute relative	Square relative	RMSE	RMSE_log	a1	a2	a3
0.066	0.416	3.359	0.147	0.945	0.974	0.986

Depth Map predictions – SPIder depth network pre-trained model



Dataset pre-procesing

KITTI Data-set reduction



75,000 images



Train split
(70,000 images)



Validation split
(5,000 images)



Detect car images using
YOLOv8

Filter the images that
contains 1-3 cars



New split files



Train split file
(2,500
images)

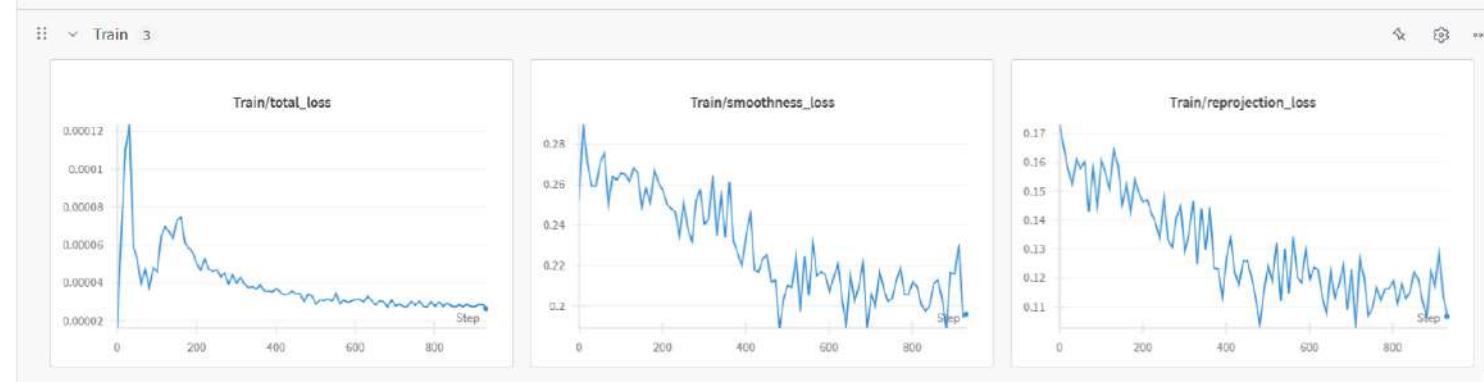


Validation
split file (200
images)

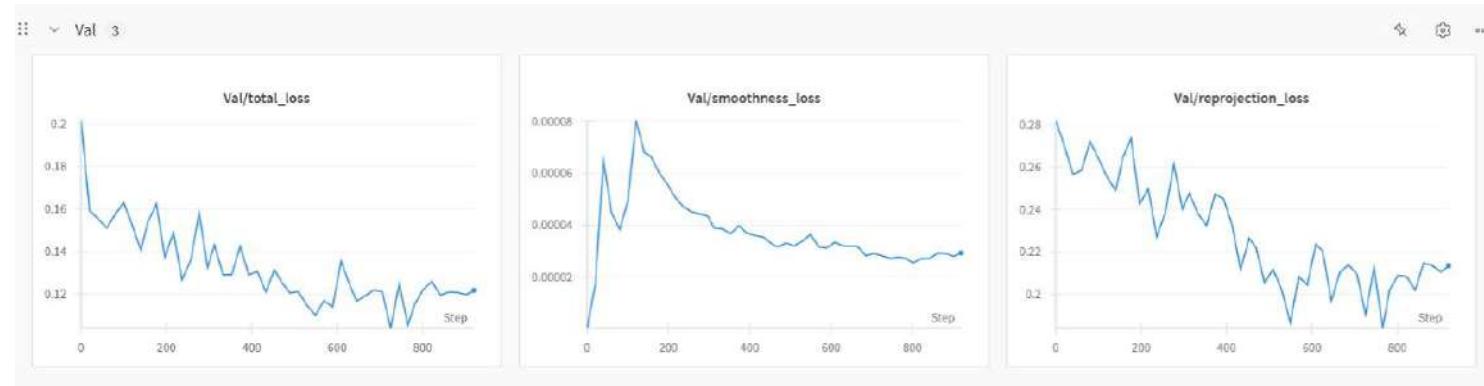
SPIdepth network – KITTI Baseline training

Hyperparameters

- Dataset: 2,500 train images
- Epochs: 6
- Batch size: 16
- Image resolution: 640x192
- Encoder backbone: ResNet-18 Lite



Converged around 0.000027 after 938 steps



Converged around 0.12 after 938 steps

Inference on Reduced validation set

Validation set of 200 images, filtered using YOLO

Absolute	Square	RMSE	RMSE_log	a1	a2	a3
relative	relative					
0.157	1.142	5.481	0.223	0.795	0.933	0.976

SPIdepth network – KITTI Curriculum Learning

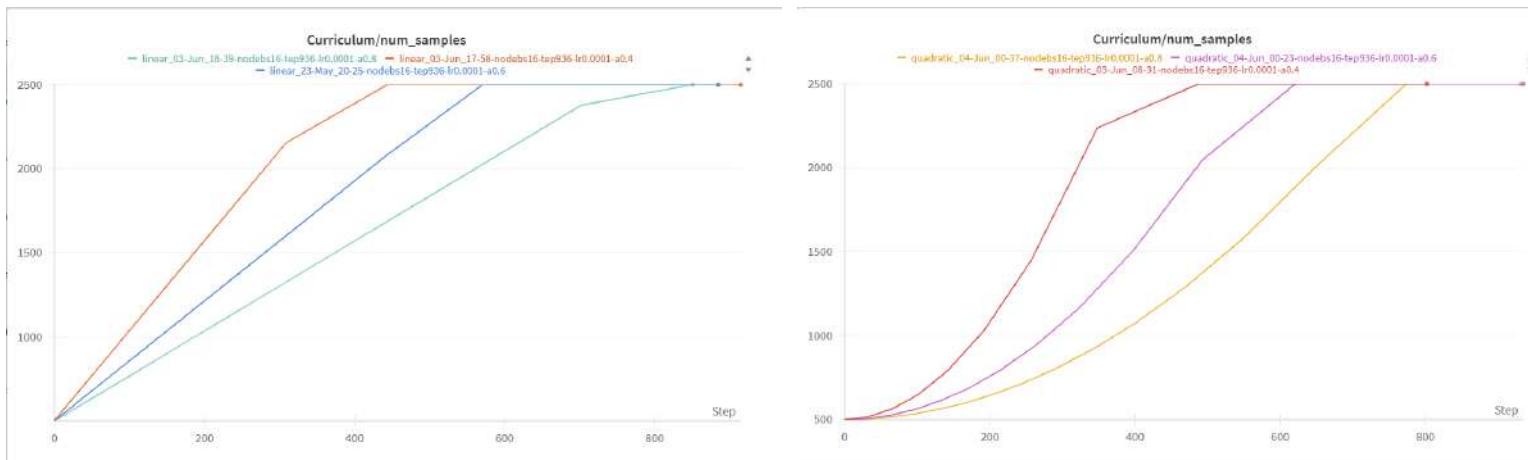
Scoring function $L = \mu L_p + \lambda L_s$

- Self-taught scoring function: Using saved weights from Baseline training (after epoch 3)
- Transfer scoring function: Pretrained SPIdepth model provided by the author



Pacing function

- T: 936 steps (6 epochs)
- N: 2500 samples
- b: 0.2 (fixed)
- a: {0.4, 0.6, 0.8} -> a controls the speed at which the full dataset is revealed during training



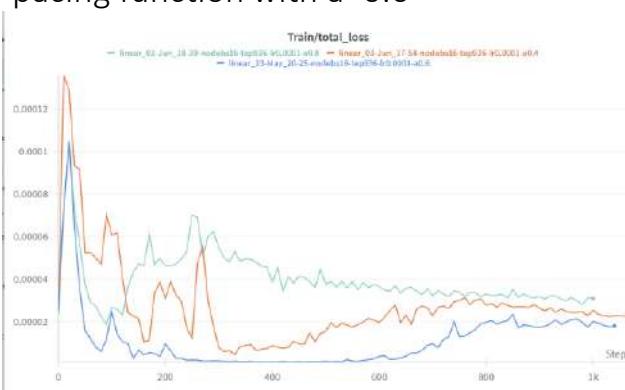
TRAINING EXPERIMENTS

10 experiments:

- **Pacing function hyperparameters (6 experiments):** Linear and Quadratic pacing function with different a values {0.4, 0.6, 0.8} and 936 steps
- **Faster convergence (3 experiments):** Linear pacing function with different a values {0.4, 0.6, 0.8} with less steps (468 steps).
- **Self-taught vs Transfer scoring (1 experiment):** Using a Linear pacing function with a=0.6

a	Absolute relative	Square relative	RMSE	RMSE_log	a1	a2	a3
0.4	0.180	1.333	6.606	0.257	0.726	0.906	0.967
0.6	0.242	2.775	10.313	0.415	0.581	0.793	0.884
0.8	0.156	1.113	5.339	0.220	0.797	0.936	0.978

Inference with Linear Pacing function with a=0.8, achieved a better result compared with Baseline training

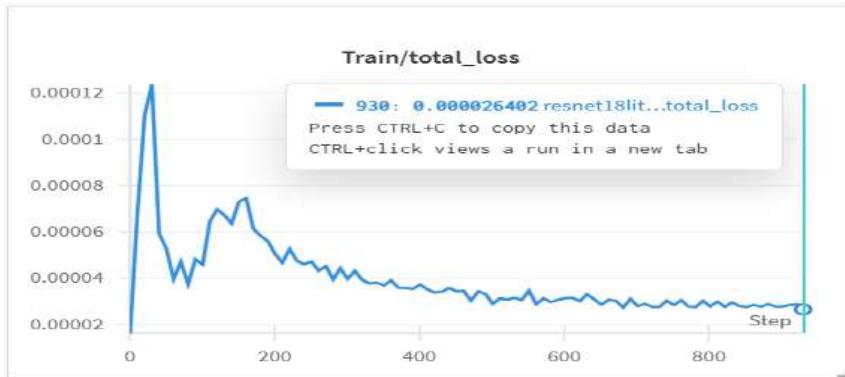


Training converged around 0.000017 (a=0.6)

SPIdeth network – Curriculum Learning

CONVERGENCE IN TRAINING

Linear pacing function ($a=0.6$) reduced the Training loss from 0.000027 to 0.000017. (Better convergence)



Baseline training



Curriculum Learning

GENERALIZATION

Linear pacing function ($a=0.8$) generalized better for unseen data.

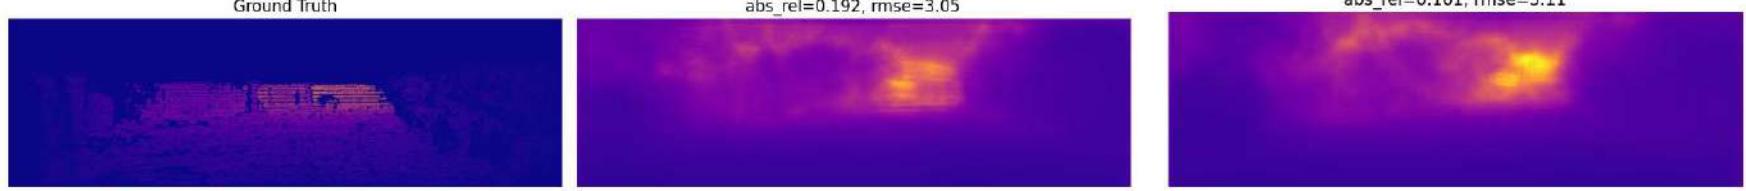
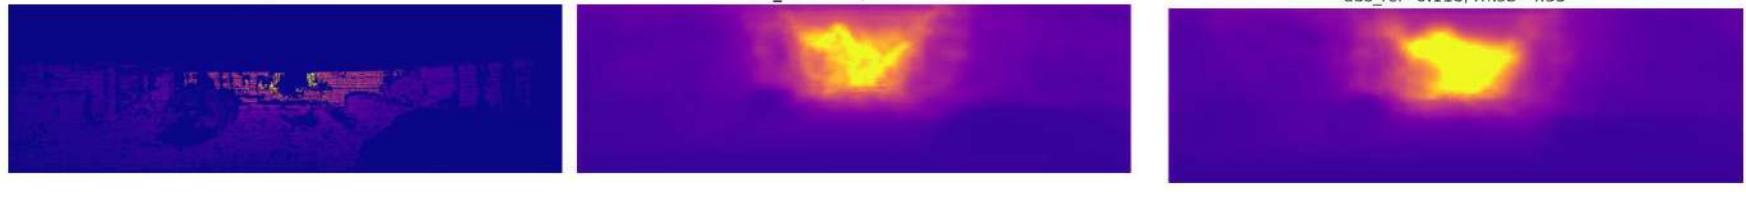
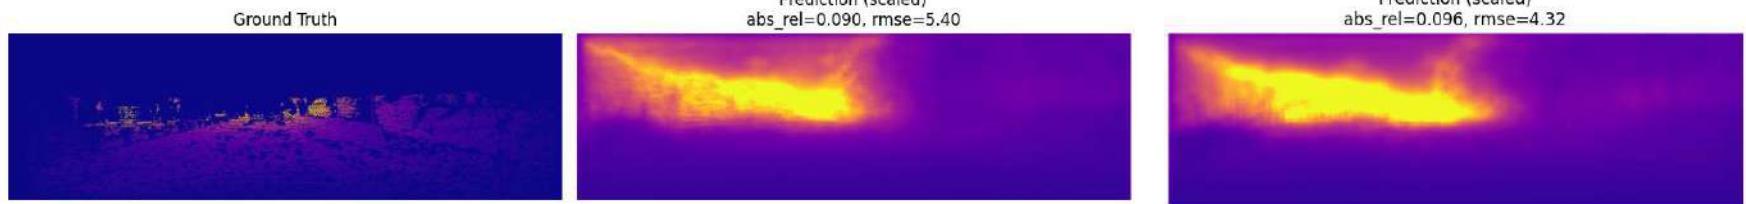
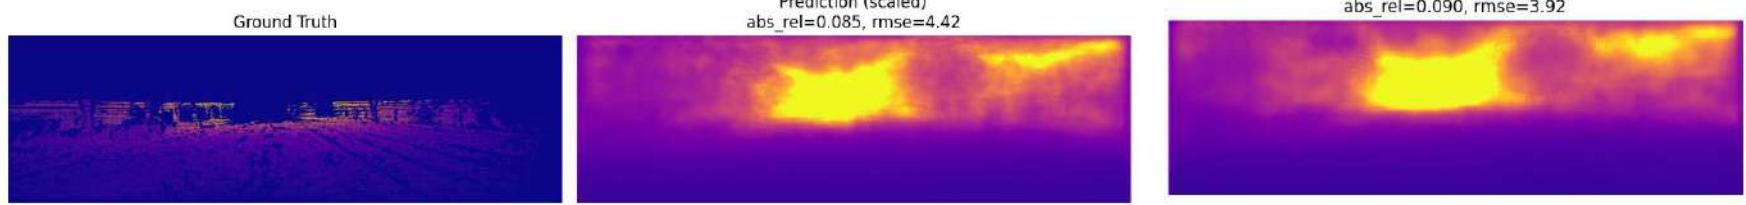
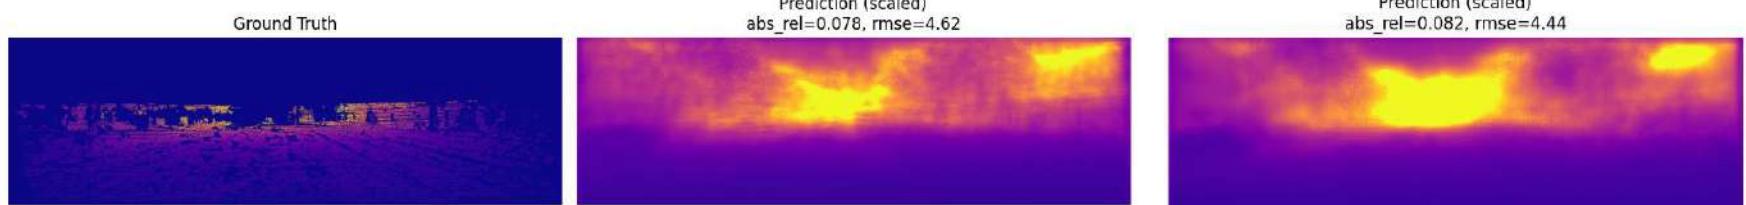
Absolute relative	Square relative	RMSE	RMSE_log	a1	a2	a3
0.157	1.142	5.481	0.223	0.795	0.933	0.976

Baseline training

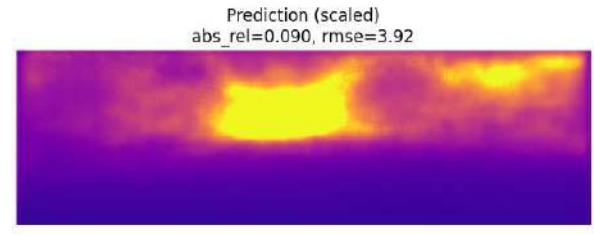
a	Absolute relative	Square relative	RMSE	RMSE_log	a1	a2	a3
0.4	0.180	1.333	6.606	0.257	0.726	0.906	0.967
0.6	0.242	2.775	10.313	0.415	0.581	0.793	0.884
0.8	0.156	1.113	5.339	0.220	0.797	0.936	0.978

Curriculum Learning

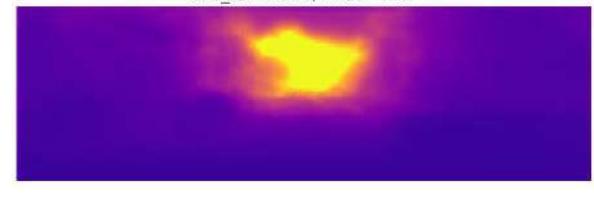
Depth Map predictions – KITTI Curriculum Learning Training (a=0.8 best generalization) vs Baseline Training



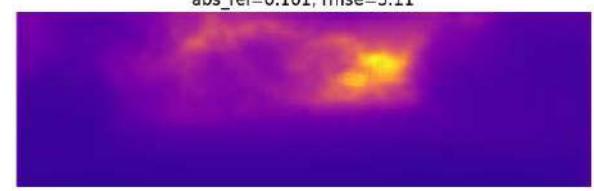
Prediction (scaled)
abs_rel=0.082, rmse=4.44



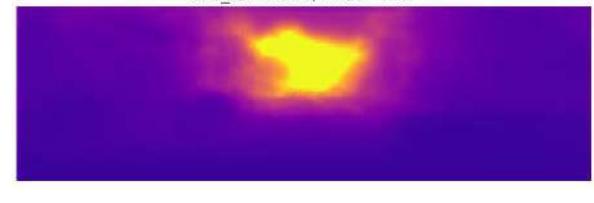
Prediction (scaled)
abs_rel=0.090, rmse=3.92



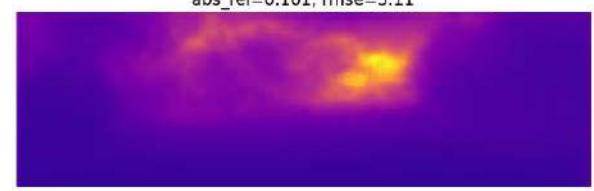
Prediction (scaled)
abs_rel=0.096, rmse=4.32



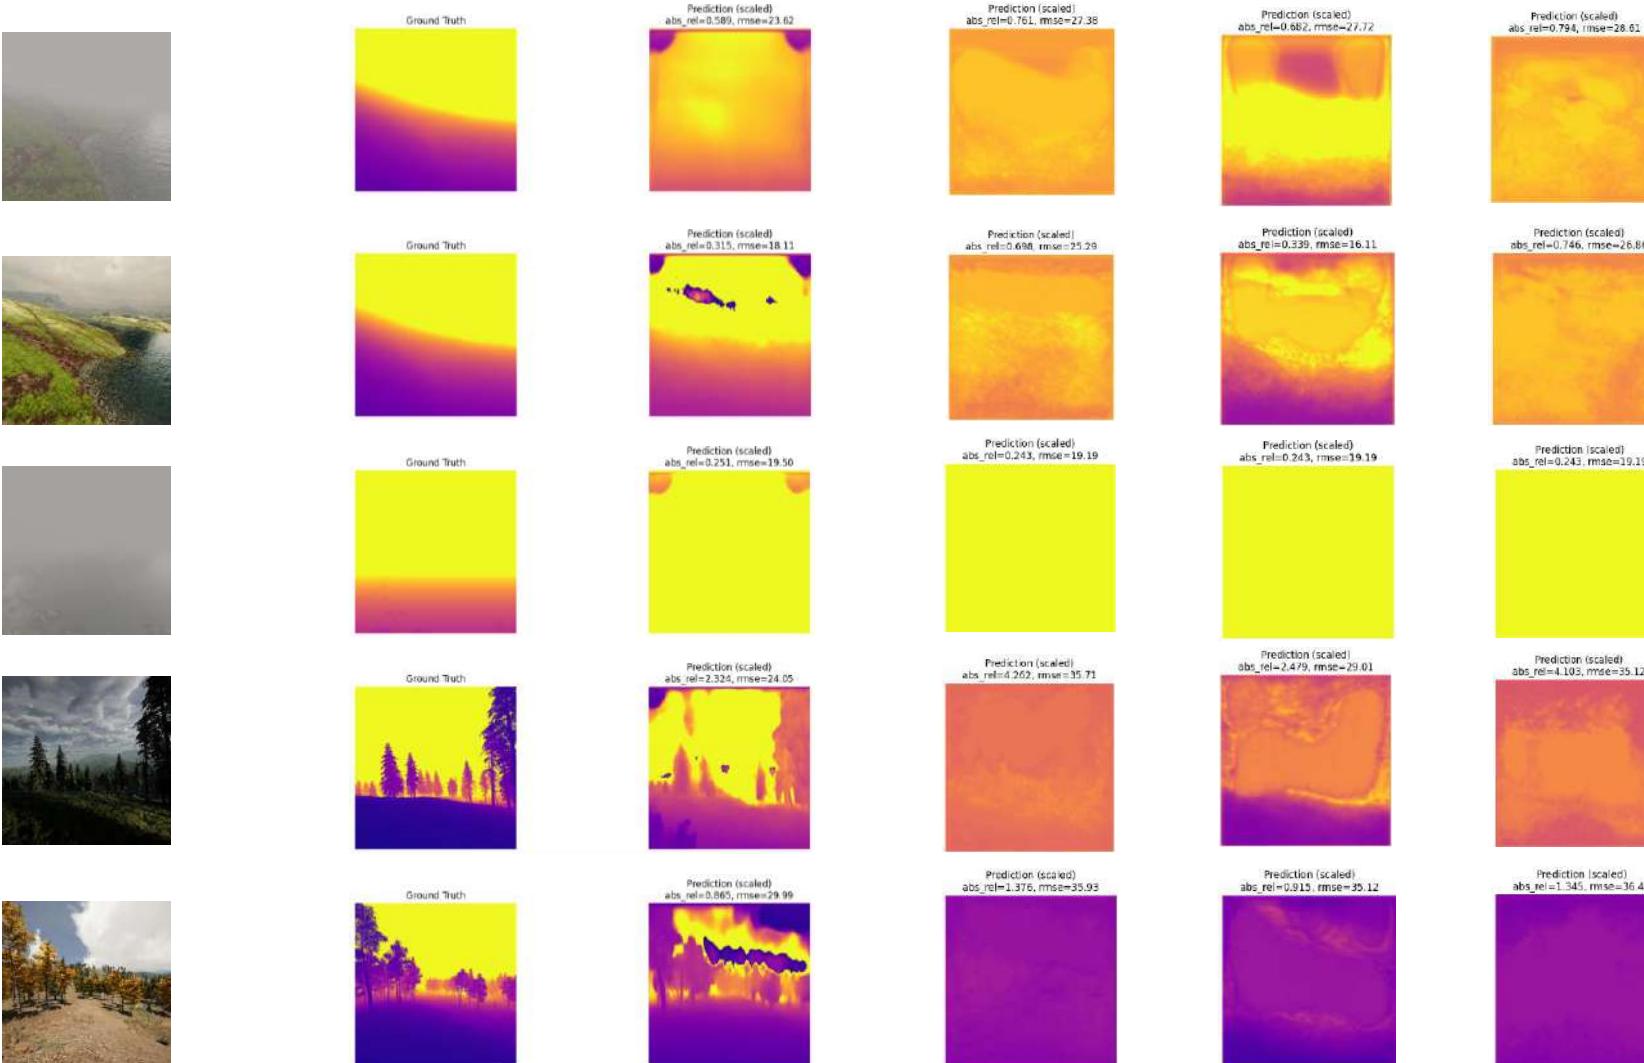
Prediction (scaled)
abs_rel=0.118, rmse=4.35



Prediction (scaled)
abs_rel=0.161, rmse=3.11



Depth map predictions MidAIR – SOTA SPIdepth vs Baseline Training ep 9 vs MidAIR Curriculum Learning Training (a=0.8 best generalization) ep 8 vs ep 9



Model	Dataset	Training	Epochs	Baseline		CL			
				Abs	Rel	RMSE	Abs	Rel	RMSE
SPIdepth	KITTI	Unsupervised	6	0.157	5.481	0.156	5.339		
	MidAir	Unsupervised	9	1.126	33.273	1.125	33.371		
Monodepthv2	KITTI	Unsupervised	–	–	–	–	–	–	–
	MidAir	Unsupervised	9	1.162	34.702	1.157	33.449		

CONCLUSIONS

- The experiments demonstrated that Curriculum Learning not only led to faster convergence but also improved generalization in early stages of training. This behavior suggests that CL can accelerate convergence during the initial iterations, which is particularly advantageous in time-constrained scenarios such as disaster response, where rapid deployment is critical.

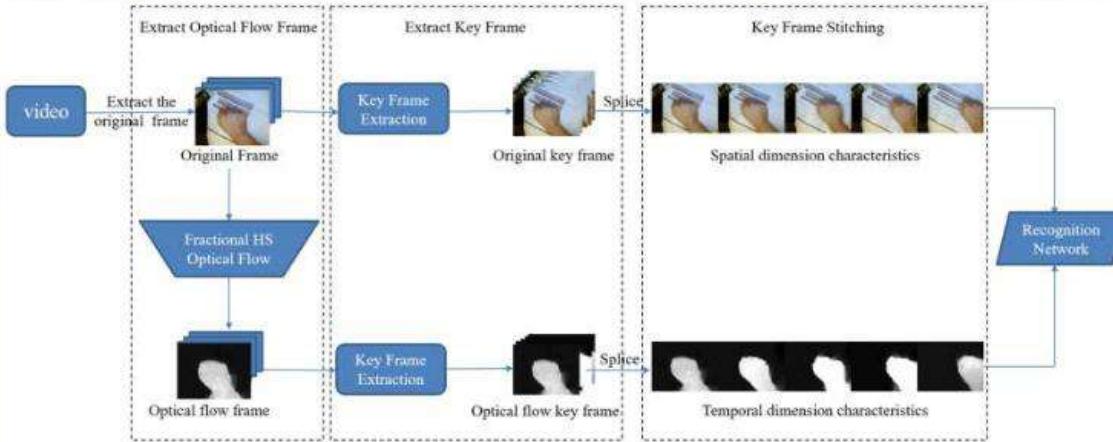
What is next?:

- Perform more experiments using Supervised Learning
- Publish these results in a journal paper

Projects

Real-time Hand Gesture Recognition for Interactive Video Conferencing

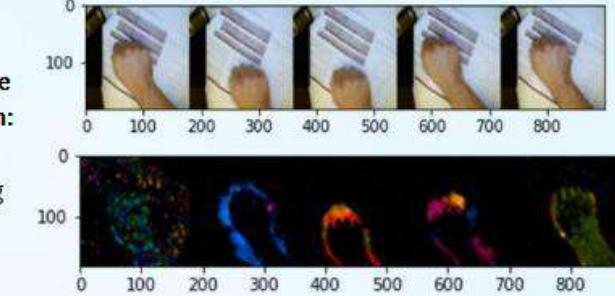
1. Machine Vision: Developing and training the CNN using Optical Flow and Key Frame Extraction techniques.



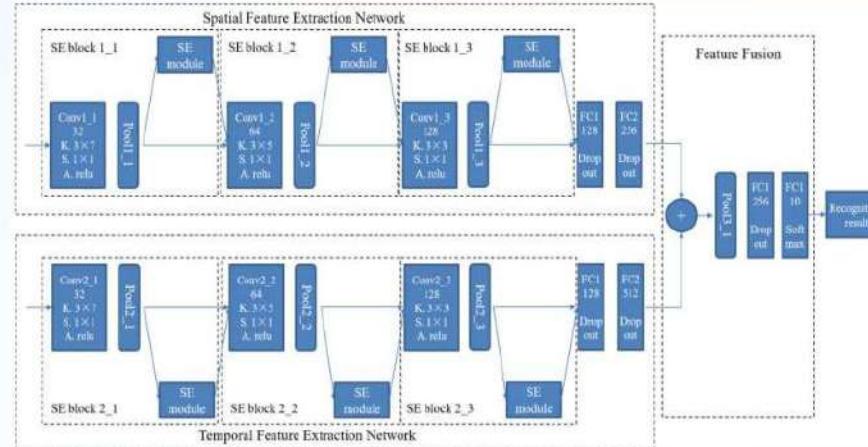
Optical Flow:
Gunnar-Farneback flow



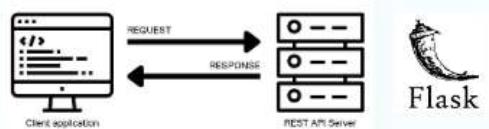
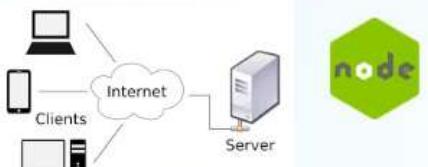
Key Frame Extraction:
5 frames
Clustering



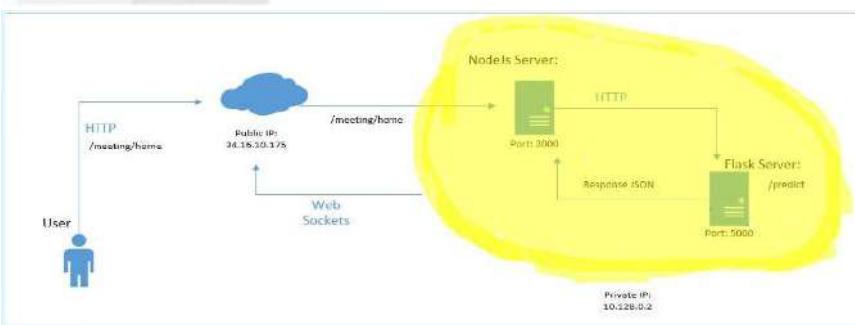
CNN architecture:



2. Web server development: Backend Server API server (Flask) and Frontend server NodeJS



3. Deployment: Hosting the application on Google Cloud Platform to ensure scalability and reliability.



Flask Server: Serves the API (/predict) endpoint at Port 5000

```
PS D:\PythonProjects\HandRecognition> python flaskServer.py
 * Serving Flask app "flaskServer"
 * Debug mode: off
WARNING: This is a development server. Do not use it in a prod
* Running on http://127.0.0.1:5000
* Running on https://127.0.0.1:5001
* Running on http://102.168.1.38:5000
Press CTRL+C to quit
```

Nodejs Server: Serves the Web Server at port 3000

```
PS D:\[Projects]Node\HandMovementRecognition\NodeWebApp> node server.js
Server running on port 3000
```



Welcome to the Video Call App

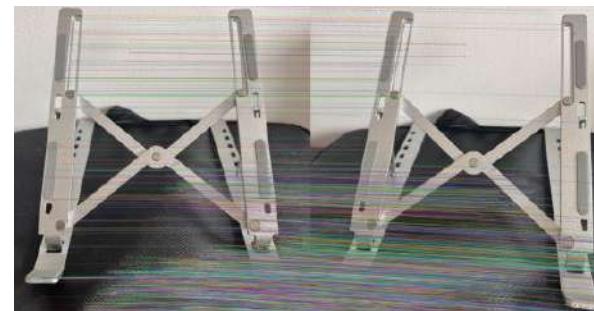
Join the Room

Projects

3D Reconstruction using Structure from Motion and NeRFs



OpenCV



Key point detection and Feature matching (SIFT and FLANN)

Triangulation

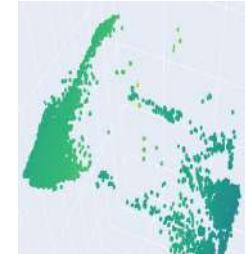
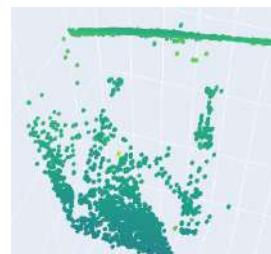
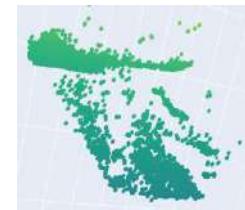
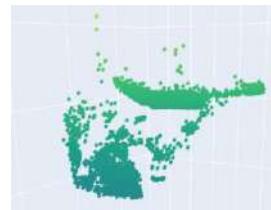


• • •

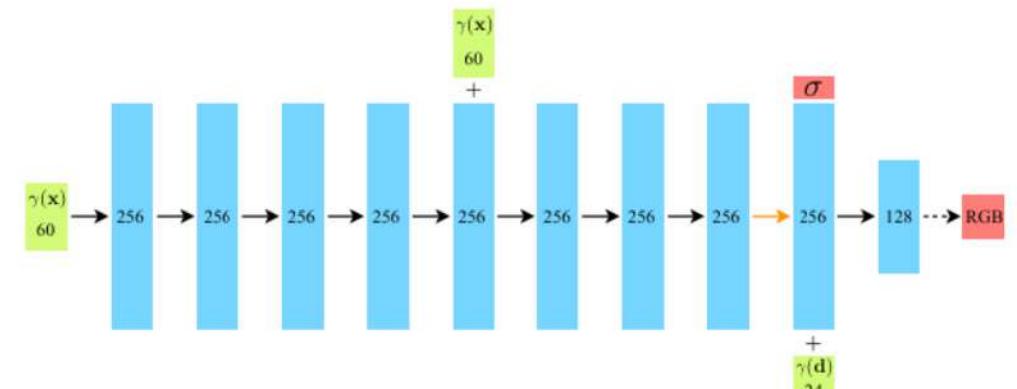


20 images (iphone camera)

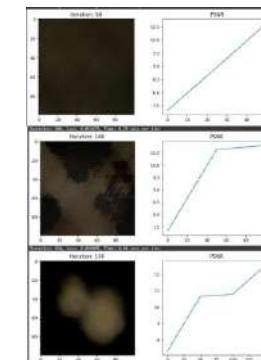
Structure from Motion



Input: Ray position and direction

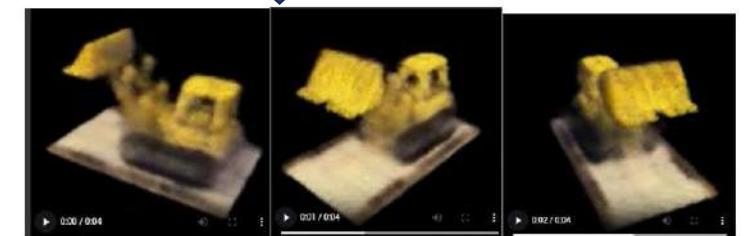


Training



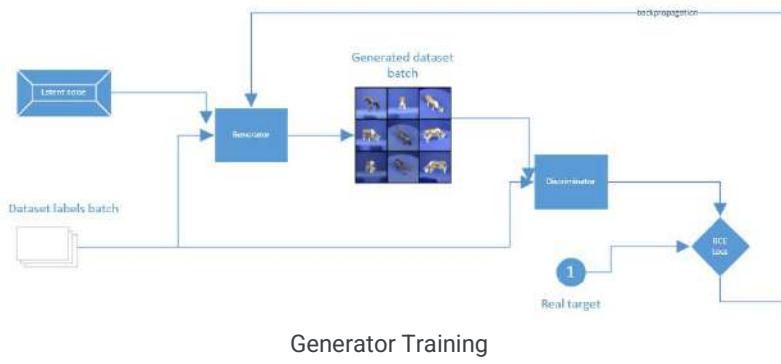
NeRF

Volume rendering

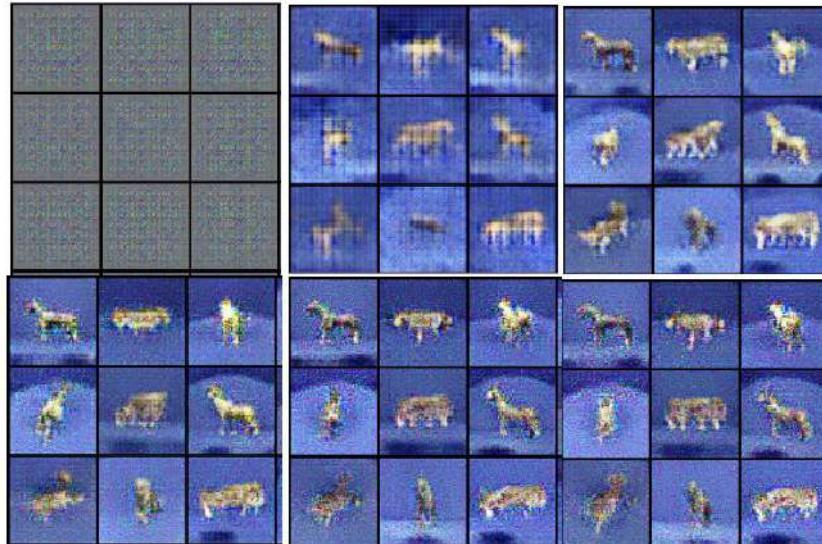


Projects

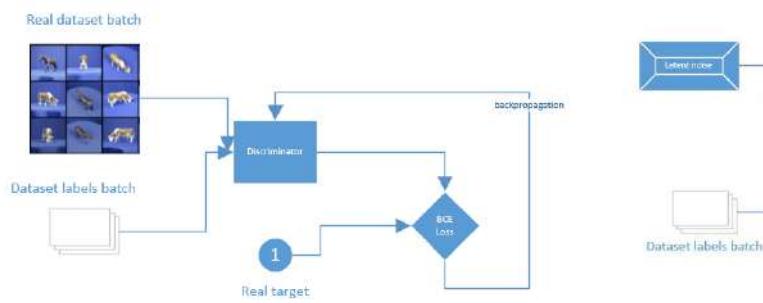
Generative Adversarial Networks for Data Augmentation



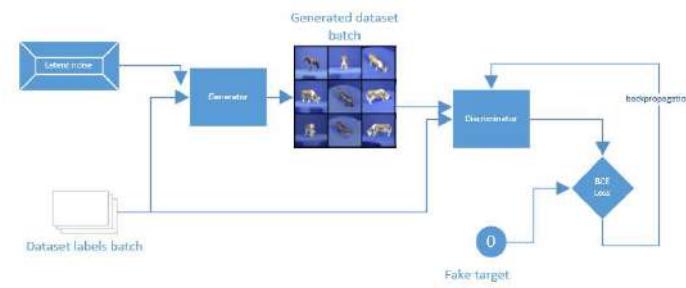
Generator Training



Training



Discriminator Training

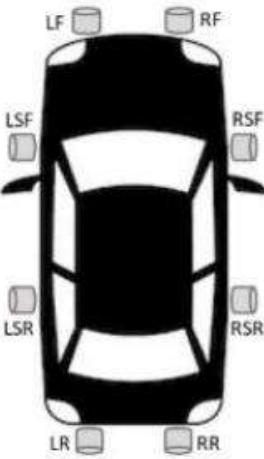


Fake Images

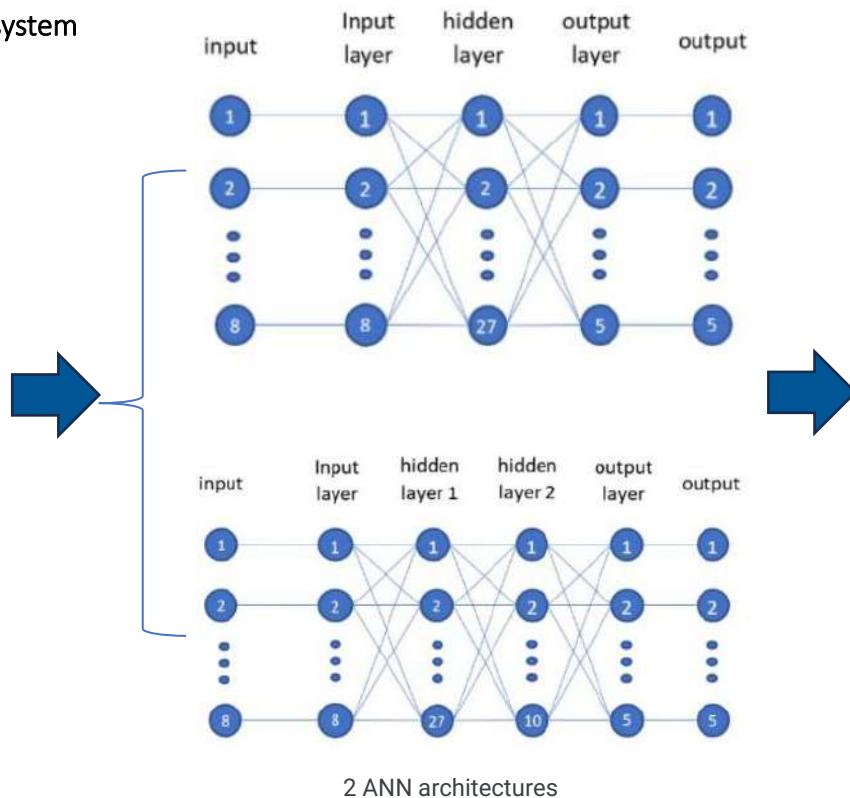


Projects

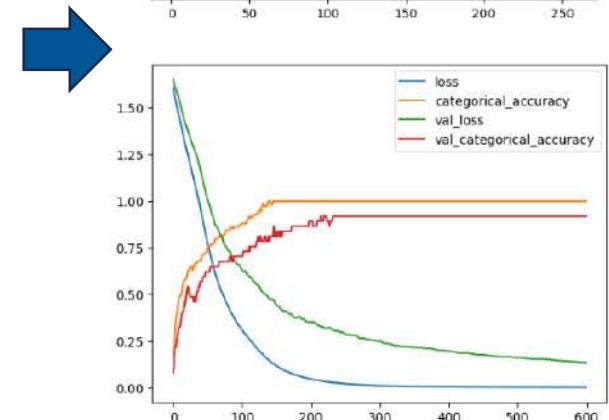
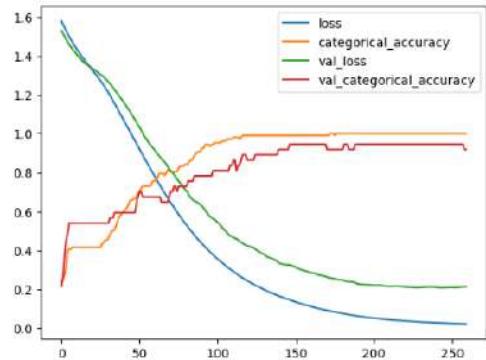
ANN for Autonomous vehicle Ultrasonic Multi-sensor system



Sensor's positions



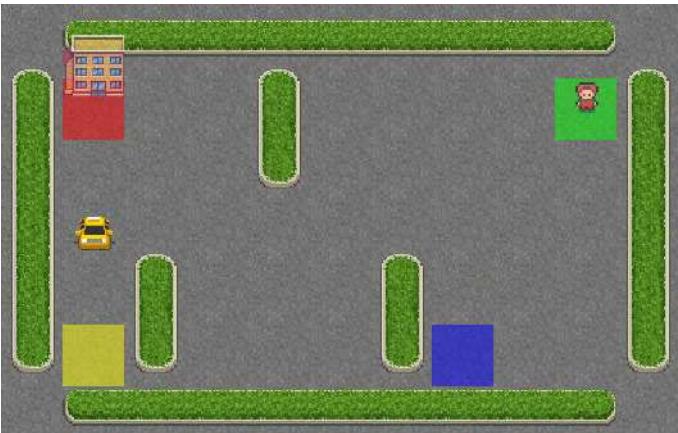
- 5 predictions:
- Stop/Braking
 - Soft Braking
 - Turn Left
 - Turn right
 - Go straight



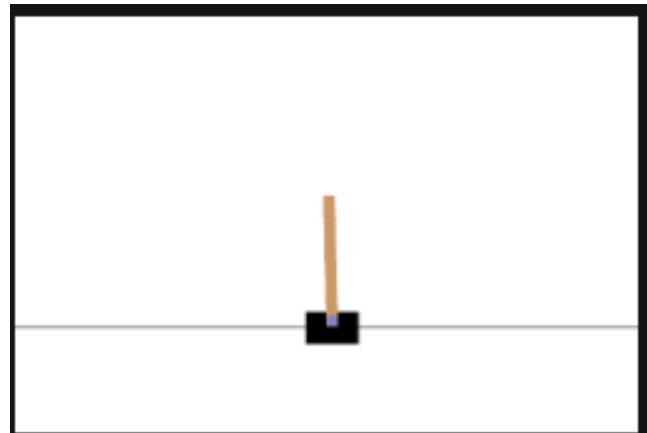
Performance of both networks

Projects

Reinforcement Learning mini projects



Q-Learning using Taxi-gym environment



DQN using Cart-Pole Gym environment



Trello Drone

- Actions:
- Go forward
 - Go back
 - Turn clockwise (90 degrees)
 - Turn anticlockwise (90 degrees)



RESNET CNN to
get the features



Target plant

```
# Open drone video stream
videoUDP = 'udp://192.168.10.1:11111'
cap = cv2.VideoCapture(videoUDP)
```

Frame captured by the drone camera



Reward: Cosine
similarity
between both:
Target features
and frame
features

Projects

Video Segmentation for a Tennis Clip

YOLO v5 to Segment
the Tennis ball

Kalman Filter using the YOLO predictions
to find the trajectory

```
# Perform YOLOv5 detection on the frame
results = model(frame)

# Process YOLOv5 detections
detected_ball = None
for result in results.xywh[0]: # Format: [x_center, y_center, width, height, confidence, class_id]
    x_center, y_center, w, h, conf, cls = result.tolist()
    if int(cls) == 32: # Replace with the class ID for 'tennis ball' (check with your model)
        detected_ball = (int(x_center), int(y_center))
```



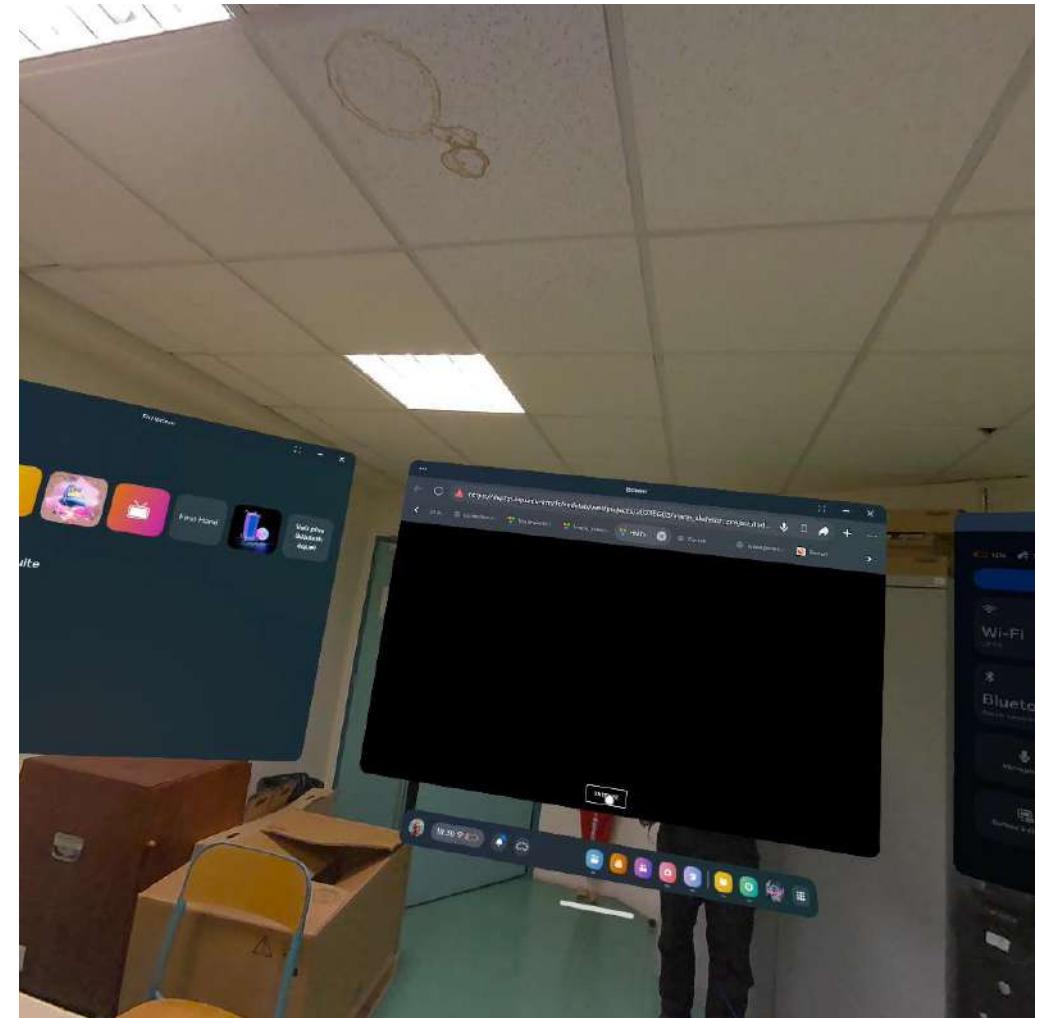
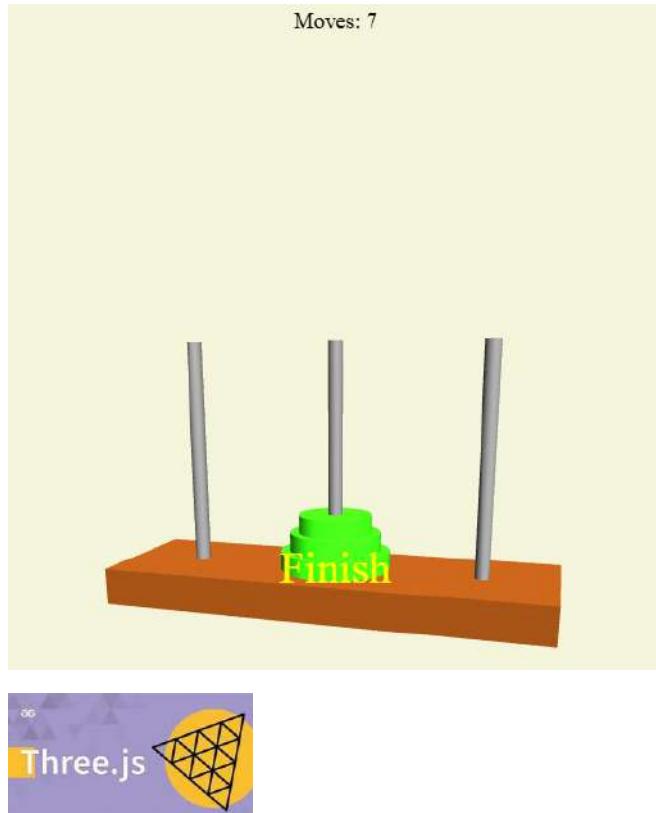
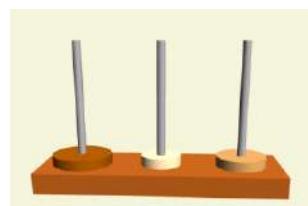
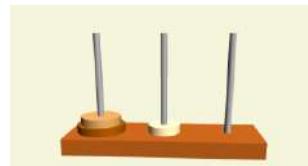
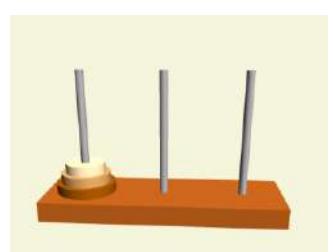
Find the speed

```
# Calculate speed (in pixels per frame)
if last_position:
    dt = 1 / fps # Time difference between frames (in seconds)
    dx = predicted_x - last_position[0]
    dy = predicted_y - last_position[1]
    speed = np.sqrt(dx**2 + dy**2) / dt # Speed in pixels per second
```



Projects

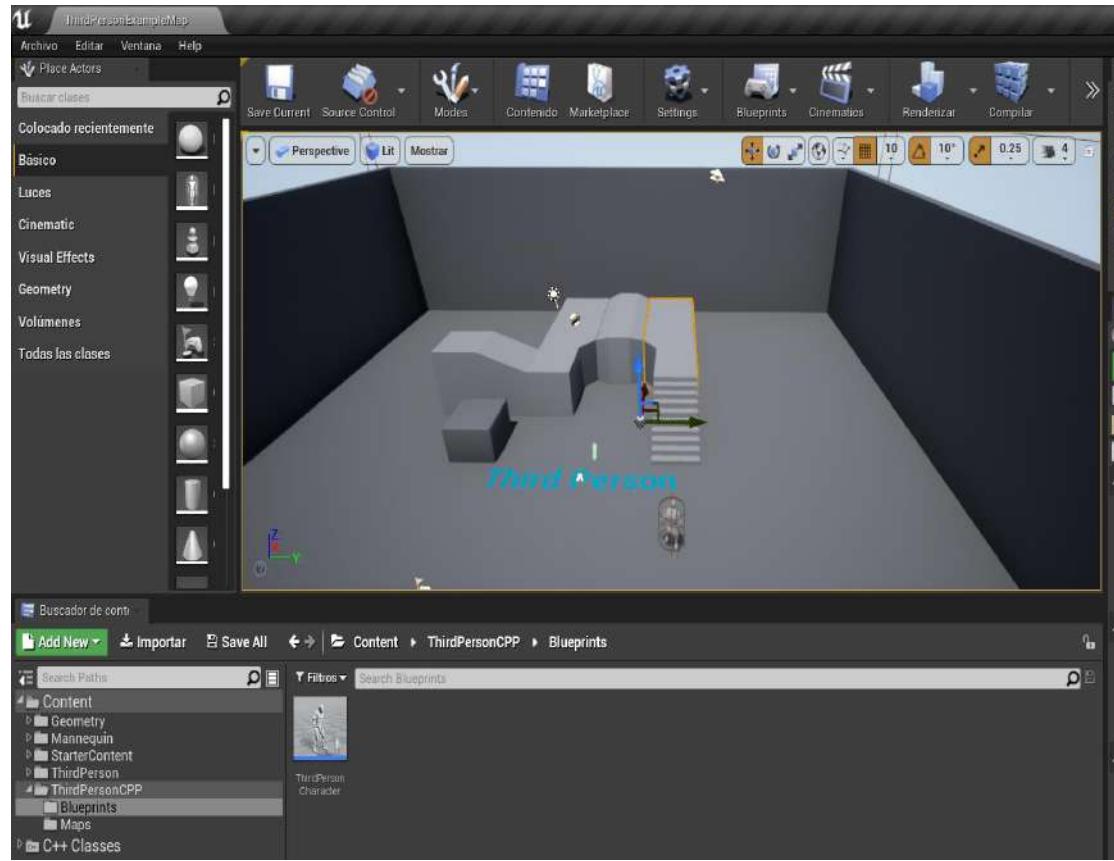
Tower of Hanoi XR implementation



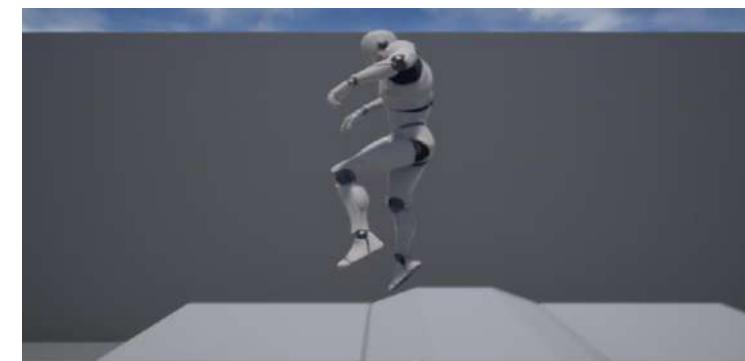
Meta Headsets

Job related experience

Basics of Unreal Engine



Using Blueprints





MERCI!